

## Task-03

### Import necessary libraries

```
library(dplyr)

##
## Attaching package: 'dplyr'
##
## The following objects are masked from 'package:stats':
##
##   filter, lag
##
## The following objects are masked from 'package:base':
##
##   intersect, setdiff, setequal, union

library(ggplot2)
library(knitr)
```

### Loading data

```
reviews <- read.csv("../data/reviews.csv", stringsAsFactors = FALSE)
```

### Identify the 10 users who have written the most reviews

Identify the 10 most active users by counting the number of reviews per *user\_id* from the reviews dataset, sorting them in descending order based on the number of reviews, and selecting the top 10 users *top\_users\_by\_count*. Next, filters the original reviews dataset to include only reviews from these top 10 users *reviews\_from\_top\_users*. After that, for each of the 10 users, the algorithm calculates the average star rating they have given, ignoring any missing values NA in the calculation *average\_stars\_for\_top\_users*.

```
top_users_by_count <- reviews %>%
  group_by(user_id) %>%
  summarise(review_count = n(), .groups = 'drop') %>% # Group the rows in the reviews dataset by the user_id
  arrange(desc(review_count)) %>% # sorting the results
  slice_head(n = 10) # get 10 rows

reviews_from_top_users <- reviews %>%
  filter(user_id %in% top_users_by_count$user_id)

# calculate average stars for each of these top users.
# na.rm = TRUE ensures that NA values in the 'stars' column are ignored during mean calculation.
average_stars_for_top_users <- reviews_from_top_users %>%
  group_by(user_id) %>%
  summarise(average_stars = mean(stars, na.rm = TRUE), .groups = 'drop')

print(average_stars_for_top_users)

## # A tibble: 10 x 2
##   user_id   average_stars
##   <chr>         <dbl>
## 1 ""             3.00
```

```
## 2 "u_11229"          3.07
## 3 "u_11551"          3.27
## 4 "u_14899"          2.57
## 5 "u_17629"          2.21
## 6 "u_22933"          2.93
## 7 "u_23971"          2.36
## 8 "u_27070"          2.83
## 9 "u_27907"          3.43
## 10 "u_6766"          3.27
```

Then, merges `left_join` the review count information `top_users_by_count` with the average star ratings `average_stars_for_top_users` for each user into a single summary table called `summary_top_users`, which is then sorted again based on the number of reviews

```
summary_top_users <- top_users_by_count %>%
  left_join(average_stars_for_top_users, by = "user_id") %>%
  arrange(desc(review_count)) # Re-arrange to ensure order by review_count

summary_table <- kable(summary_top_users,
  caption = "Top 10 Users: Review Count and Average Stars",
  col.names = c("User ID", "Total Reviews", "Average Stars"),
  align = "c",      # Center align columns
  digits = 2,       # Round average stars to 2 decimal places
  format = "pipe") # Use "pipe"

print(summary_table)
```

```
##
##
## Table: Top 10 Users: Review Count and Average Stars
##
## | User ID | Total Reviews | Average Stars |
## | :-----: | :-----: | :-----: |
## |      |      5829      |      3.00      |
## | u_27070 |      18      |      2.83      |
## | u_11551 |      15      |      3.27      |
## | u_6766  |      15      |      3.27      |
## | u_11229 |      14      |      3.07      |
## | u_14899 |      14      |      2.57      |
## | u_17629 |      14      |      2.21      |
## | u_22933 |      14      |      2.93      |
## | u_23971 |      14      |      2.36      |
## | u_27907 |      14      |      3.43      |
```

## Visualize

```
# Ensure that the order of users in the plot is consistent with the previous summary table
reviews_from_top_users$user_id <- factor(reviews_from_top_users$user_id,
  levels = summary_top_users$user_id)


# Create the boxplot
# na.rm = TRUE in geom_boxplot will ensure NA star values are ignored for plotting
rating_distribution_plot <- ggplot(reviews_from_top_users,
  aes(x = user_id,
    y = stars,
    fill = user_id)) +
```

```

geom_boxplot(na.rm = TRUE) +
labs(title = "Star Rating Distribution for Top 10 Users",
      x = "User ID",
      y = "Stars Given") +
theme_minimal(base_size = 12) +
theme(axis.text.x = element_text(angle = 45, hjust = 1, vjust = 1), # Rotate x-axis labels for readability
      legend.position = "none") + # Hide legend
coord_cartesian(ylim = c(0, 6))

print(rating_distribution_plot)

```



The boxplot shows how the top users vary in the way they give star ratings. Users such as u\_6766 and u\_27907 tend to give higher ratings overall, with medians above 3 and upper quartiles close to 5. In contrast, users like u\_17629 and u\_22933 give lower ratings more frequently, as indicated by their lower medians and compressed upper ranges. Some users show wide variability (e.g., u\_6766, u\_27907), while others have more consistent rating patterns (e.g., u\_14899). Outliers in a few users suggest occasional extreme ratings.